

Contents

Chapter 18: Agentic Development Loops	1
1. From Code Generation to Iterative Engineering	3
2. The Loop as a Control System	4
3. The Seven Stages	5
4. Stage 1 — SPEC	5
5. Stage 2 — PLAN	6
6. Stage 3 — IMPLEMENT	7
7. Stage 4 — TEST	7
8. Stage 5 — VERIFY	8
9. Stage 6 — FIX	9
10. Stage 7 — RETEST	10
11. Verifiers as Sensors	11
12. The Verifier Stack	11
13. Why Feedback Quality Matters	13
14. The Difference Between Detection and Diagnosis	13
15. The Loop as Search	14
16. Why Better Feedback Can Beat a Better Model	15
17. Iteration Depth	16
18. Stopping Conditions	17
19. Progress Measurement	17
20. Regression Control	18
21. Verification Should Be Independent	19
22. LLM-as-Judge	20
23. Benchmarks as Verifiers	21
24. Security Scanners as Verifiers	21
25. Browser Tests and Simulators	22
26. The Autonomous Development Loop	23
27. The Harness Must Control the Loop	24
28. Autonomous Does Not Mean Unbounded	24
29. The Development Loop as a Learning System	25
30. The Most Important Design Principle	26
31. Exercise — Build an Autonomous Development Loop	27
32. A More Advanced Exercise — Improve the Loop Without Chang- ing the Model	28
33. Key Takeaways	30

Chapter 18: Agentic Development Loops

The previous chapters established three foundations of agentic software engineering:

- **Chapter 15:** coding agents operate through a closed loop of reasoning, tool use, execution, and verification.
- **Chapter 16:** specifications define what the agent is supposed to produce.

- **Chapter 17:** context engineering determines what information the agent needs in order to reason effectively.

Chapter 18 brings these ideas together.

The central abstraction is the **agentic development loop**:

```

SPEC
↓
PLAN
↓
IMPLEMENT
↓
TEST
↓
VERIFY
↓
FIX
↓
RETEST
^ (loop back)

```

This looks superficially similar to traditional software development.

The difference is that the loop can now be **executed continuously by an AI system**.

Instead of:

```

Human writes code
      ↓
Human runs tests
      ↓
Human diagnoses failure
      ↓
Human fixes code

```

we can construct:

```

Agent writes code
      ↓
Verifier runs
      ↓
Agent observes failure
      ↓
Agent diagnoses
      ↓
Agent fixes code
      ↓
Verifier runs again

```

This produces a profound shift in the engineering problem.

The objective is no longer simply to make the model generate better code.

It is:

Design a development loop in which the system can detect its own mistakes and progressively converge toward a verified solution.

1. From Code Generation to Iterative Engineering

A traditional LLM coding interaction often looks like:

Prompt

↓

LLM

↓

Code

The problem is that the model has to be correct immediately.

An agentic system instead looks like:

Task

↓

LLM

↓

Implementation

↓

Verifier

↓

Feedback

↓

LLM

↓

Correction

↓

Verifier

~ (loop back)

The model no longer needs to solve the entire problem in one inference.

It needs to:

1. make a reasonable decision,
2. observe its consequences,
3. interpret the evidence,
4. update its approach,

5. try again.

This is much closer to how experienced engineers actually work.

A developer rarely writes a complex feature perfectly on the first attempt.

Instead:

```
implement
→ compile
→ test
→ inspect
→ modify
→ test
→ debug
→ test
→ review
```

Agentic development operationalizes this process.

2. The Loop as a Control System

The development loop can be formalized as a feedback-control process.

Let:

$$S_t$$

represent the current state of the software system.

The specification defines the desired set of states:

$$\mathcal{G} = S \mid S \models SPEC$$

The agent chooses an action:

$$A_t \sim \pi_\theta(A \mid C_t, S_t)$$

which changes the system:

$$S_{t+1} = T(S_t, A_t)$$

A verifier evaluates the new state:

$$V(S_{t+1}, SPEC) \rightarrow F_{t+1}$$

where F is feedback.

The agent then uses that feedback to select the next action.

Thus:

$$S_t \rightarrow A_t \rightarrow S_{t+1} \rightarrow V \rightarrow F \rightarrow A_{t+1}$$

The development process becomes a **closed-loop optimization problem**.

The goal is to reach:

$$S_n \in \mathcal{G}$$

with sufficiently high confidence.

3. The Seven Stages

The canonical loop is:

SPEC
↓
PLAN
↓
IMPLEMENT
↓
TEST
↓
VERIFY
↓
FIX
↓
RETEST

Each stage has a distinct role.

4. Stage 1 — SPEC

The loop begins with a specification.

For example:

Build a REST API for document search.

Requirements:

- `POST /query`
- `authenticated users only`
- `tenant isolation`
- `citations required`
- `P95 latency < 500 ms`

Acceptance:

- `all tests pass`
- `unauthorized requests return 401`
- `cross-tenant retrieval is impossible`
- `latency target is satisfied`

The specification establishes the target.

Without it, the agent cannot reliably determine whether its work is complete.

This is why Chapter 16 matters.

The agentic loop is only as good as the target against which it is evaluated.

5. Stage 2 — PLAN

The agent determines how to reach the desired state.

A plan might be:

1. `Inspect existing API structure.`
2. `Locate authentication middleware.`
3. `Inspect retrieval service.`
4. `Add /query endpoint.`
5. `Add tenant filtering.`
6. `Add citation generation.`
7. `Add tests.`
8. `Run test suite.`
9. `Benchmark latency.`
10. `Fix failures.`

The plan is not necessarily fixed.

As the agent discovers new information, it should be able to revise it.

This distinction is important:

Static plan:

`Plan once → execute blindly`

Adaptive plan:

`Plan → execute → observe → replan`

Agentic development generally benefits from the second approach.

6. Stage 3 — IMPLEMENT

The agent now modifies the repository.

This may involve:

```
read files
search symbols
edit source
create files
modify configuration
update dependencies
write tests
```

Implementation is not necessarily one action.

It may itself contain a loop:

```
inspect
↓
modify
↓
inspect
↓
modify
```

The agent should generally make changes in manageable increments.

Large uncontrolled modifications make verification and debugging harder.

7. Stage 4 — TEST

The agent executes tests.

Examples include:

```
unit tests
integration tests
end-to-end tests
regression tests
property tests
```

A good development loop does not wait until the entire feature is implemented before testing.

Instead:

small change
↓
test
↓
next change
↓
test

This reduces the debugging search space.

If a test fails immediately after a small change, the causal relationship is easier to identify.

8. Stage 5 — VERIFY

Testing is only one form of verification.

Verification asks:

Does the resulting system satisfy the specification?

Different verifiers provide different evidence.

For example:

Compiler
↓
Syntax / compilation correctness

Type checker
↓
Type consistency

Unit tests
↓
Local behavioral correctness

Integration tests
↓
Component interaction

Benchmark
↓
Performance

Security scanner
↓
Security properties

Evaluator
↓
Higher-level behavior

Human review
↓
Intent / architecture / risk
This leads to an important distinction:

$\text{Test} \subset \text{Verification}$

Testing is one verification mechanism, not the entire verification system.

9. Stage 6 — FIX

If verification fails, the agent receives evidence.

For example:

```
pytest  
FAILED tests/test_auth.py::test_cross_tenant_access
```

Expected:
403

Received:
200

The agent should now form a hypothesis:

Tenant authorization is being applied after retrieval rather than before it.

It might inspect:

```
authorization middleware  
retrieval query  
tenant filtering
```

and make a correction.

This stage is where the feedback loop becomes more powerful than one-shot generation.

The agent does not need to know the solution beforehand.

It needs the ability to **learn from the verifier's feedback within the current task trajectory**.

10. Stage 7 — RETEST

The agent reruns verification.

There are three possible outcomes:

PASS

↓

Continue to next verifier

FAIL

↓

Analyze and fix

INCONCLUSIVE

↓

Acquire more information

This last category is important.

Not every failure tells the agent exactly what is wrong.

For example:

Timeout

might indicate:

- performance problem
- network failure
- deadlock
- overloaded environment
- flaky test
- dependency failure

The agent may need additional diagnostics before attempting a fix.

Thus the loop is not simply:

failure → patch

It is:

failure

↓

diagnosis

↓

hypothesis

↓
additional evidence
↓
patch
↓
verification

11. Verifiers as Sensors

A useful systems perspective is to think of verifiers as **sensors**.

The agent cannot directly observe every property of the software.

Verifiers expose particular dimensions of system state.

For example:

Compiler
→ compilation state

Type checker
→ type consistency

Unit tests
→ local behavior

Integration tests
→ system interaction

Benchmark
→ performance

Security scanner
→ vulnerability indicators


Evaluator
→ task-level quality

No single verifier observes everything.

Therefore, a robust agentic development system uses **multiple independent signals**.

12. The Verifier Stack

A mature verification pipeline might look like:



Different layers catch different classes of errors.

For example:

Syntax error

→ compiler

Incorrect type

→ type checker

Wrong local behavior

→ unit test

Broken component interaction

→ integration test

Performance regression

→ benchmark

Security vulnerability

→ security scanner

Poor semantic answer

→ evaluator

The result is a **multi-dimensional feedback system**.

13. Why Feedback Quality Matters

Consider two verifiers.

Weak verifier

FAILED

Strong verifier

test_query_respects_tenant_boundary FAILED

Expected:

tenant_id = "A"

Actual:

tenant_id = "B"

The query returned a document belonging
to another tenant.

The second provides substantially more useful information.

The agent can immediately formulate a hypothesis.

This leads to a crucial principle:

**A verifier is valuable not only because it detects errors,
but because it provides actionable information about those
errors.**

A good verifier therefore has high **diagnostic value**.

14. The Difference Between Detection and Diagnosis

An agentic system needs more than error detection.

Suppose:

Test failed.

This detects an error.

But:

Expected 10 results.

Received 0.

Failure occurs only when filters are combined.

helps diagnose it.

The distinction is:

$$\text{Detection} \neq \text{Diagnosis}$$

A strong development loop tries to maximize both.

The ideal feedback is:

What failed?

Where did it fail?

Under what conditions?

What was expected?

What actually happened?

What evidence is available?

This dramatically reduces the search space for the agent's next action.

15. The Loop as Search

Agentic development can also be viewed as search.

Suppose the current implementation is:

$$I_0$$

The agent generates a modification:

$$I_1 = A(I_0)$$

The verifier produces:

$$V(I_1) = F_1$$

The agent then generates:

$$I_2 = A(I_1, F_1)$$

and continues:

$$I_0 \rightarrow I_1 \rightarrow I_2 \rightarrow \dots \rightarrow I_n$$

until:

$$V(I_n) = PASS$$

The verifier is therefore a **search oracle** that eliminates incorrect candidate states.

This is one reason automated verification is so powerful.

It allows the agent to explore a solution space without requiring perfect reasoning at every step.

16. Why Better Feedback Can Beat a Better Model

Consider two systems.

System A

Excellent model
+
weak tests

System B

slightly weaker model
+
excellent specification
+
strong tests
+
type checker
+
benchmark
+
security scanner

System B may outperform System A.

Why?

Because System B has a better mechanism for correcting errors.

This leads to the central principle of Chapter 18:

Don't make the model smarter. Make the feedback loop better.

This does not mean model capability is irrelevant.

It means model capability is only one term in the overall system.

A useful abstraction is:

$$Q_{\text{system}} = f(Q_{\text{model}}, Q_{\text{spec}}, Q_{\text{context}}, Q_{\text{tools}}, Q_{\text{verifiers}}, Q_{\text{recovery}})$$

Improving any of these can improve the final system.

17. Iteration Depth

A key variable is the number of iterations the agent can perform.

Let:

$$p = P(\text{successful iteration})$$

If iterations were independent—which they are not in practice—the probability of at least one success after n attempts would be:

$$1 - (1 - p)^n$$

The equation is only illustrative because real agent iterations are correlated.

Nevertheless, it captures an important intuition:

An agent that can repeatedly detect and repair errors can be much more capable than a system restricted to one attempt.

But unlimited iteration is not automatically good.

It creates risks:

- cost explosion
- infinite loops
- repeated mistakes
- oscillating fixes
- regression
- destructive changes
- diminishing returns

Therefore the loop needs **stopping conditions**.

18. Stopping Conditions

A robust agent should stop when:

Success

All acceptance criteria satisfied.

Budget exhausted

Maximum iterations reached.

Maximum tokens reached.

Maximum cost reached.

Progress stalls

No improvement across N iterations.

Error is unrecoverable

Missing dependency

Unavailable external service

Insufficient permissions

Contradictory specification

Human intervention required

Production deployment

Destructive operation

Security-sensitive decision

Ambiguous requirement

A useful policy is:

$$Stop = Success \vee BudgetExceeded \vee NoProgress \vee HumanRequired$$

Stopping conditions are part of agent engineering, not an afterthought.

19. Progress Measurement

The agent should ideally measure whether it is actually improving.

Suppose a verifier returns a score:

$$V_t$$

Then:

$$\Delta V_t = V_t - V_{t-1}$$

can indicate progress.

For example:

```
Iteration 1: 8 failing tests
Iteration 2: 5 failing tests
Iteration 3: 2 failing tests
Iteration 4: 0 failing tests
```

This is obvious progress.

But consider:

```
Iteration 1: 8 failures
Iteration 2: 6 failures
Iteration 3: 9 failures
Iteration 4: 7 failures
Iteration 5: 8 failures
```

The agent may be oscillating.

A harness can detect this and intervene.

Progress metrics can include:

```
tests passing
requirements satisfied
benchmark score
security findings
type errors
lint errors
```

20. Regression Control

One of the dangers of autonomous repair is that fixing one problem can create another.

For example:

```
Before:
95 tests passing
```

```
Fix bug A
```

After:
 98 tests passing
 This looks good.
 But perhaps:
 3 previously passing tests now fail.
 The agent has traded one problem for another.
 Therefore, the verifier should distinguish:
 newly passing tests
 newly failing tests
 unchanged tests
 A stronger acceptance criterion might be:

$$\text{Pass}_{\text{after}} \supseteq \text{Pass}_{\text{before}}$$

for regression-sensitive systems.

This creates an important principle:

An autonomous repair loop must optimize for net improvement, not merely local improvement.

21. Verification Should Be Independent

A particularly important architectural principle is **independence**.

If the same model writes the implementation and determines whether the implementation is correct, the system risks correlated errors.

For example:

LLM:
 "I believe the implementation is correct."

LLM:
 "Yes, it passes my evaluation."

This is weak evidence.

Prefer independent mechanisms:

LLM
 ↓
 Implementation
 ↓

pytest
↓
compiler
↓
type checker
↓
security scanner
↓
benchmark

Deterministic verifiers are especially valuable because they do not share the model's reasoning process.

22. LLM-as-Judge

Sometimes deterministic verification is impossible.

For example:

Is the generated answer sufficiently complete?

A model-based evaluator may be useful:

Candidate answer
↓
Evaluator model
↓
rubric
↓
score

This can work well, but introduces another probabilistic component.

Therefore:

LLM Judge $\neq$ ground truth

It should ideally be combined with:

- deterministic tests
- golden examples
- human evaluation
- structured rubrics
- multiple evaluators

The general principle remains:

Use the strongest independent evidence available.

23. Benchmarks as Verifiers

Performance requirements require different feedback.

Suppose the specification says:

$$P95 < 500ms$$

Functional tests might all pass while the system violates this requirement.

A benchmark provides a different signal:

```
Functional tests
    PASS
```

```
P95 latency
    FAIL 723 ms
```

The agent now has a concrete optimization target.

It might inspect:

```
database queries
retrieval
serialization
network calls
caching
```

and iterate.

This demonstrates why verification must correspond directly to the specification.

If the specification contains a requirement, the verification system should ideally have a mechanism capable of measuring it.

24. Security Scanners as Verifiers

Security is another dimension that ordinary tests may miss.


A feature may satisfy:

```
functional tests PASS
```

while introducing:

```
SQL injection
secret exposure
authorization bypass
unsafe deserialization
```

A security scanner adds another feedback channel:




This makes security analysis part of the development loop rather than a final manual audit.

25. Browser Tests and Simulators

For systems with external behavior, verification may involve simulated environments.

Examples:



The general pattern is unchanged:


Implementation → Environment → Observation → Feedback

The verifier does not have to be a test runner.


It is any mechanism that produces evidence about whether the system satisfies its specification.


26. The Autonomous Development Loop

A practical coding agent can therefore implement:



A more realistic implementation adds:





This is the architecture of an autonomous software-development loop.

27. The Harness Must Control the Loop

The model should not be solely responsible for deciding when the loop ends.

The harness should enforce:

```
maximum iterations
maximum cost
maximum execution time
allowed tools
allowed directories
verification requirements
human approval thresholds
```

For example:

```
if tests_pass
    and security_scan_pass
    and benchmark_pass:
    accept

elif iteration_count >= 20:
    request_human

elif no_progress >= 3:
    request_human

else:
    continue
```


The agent supplies reasoning.

The harness supplies **control policy**.

28. Autonomous Does Not Mean Unbounded

An autonomous loop should have bounded autonomy.

A useful architecture is:



This is especially important for:

- production infrastructure
- databases
- security configuration
- destructive operations
- external communication
- financial systems

The goal is not maximum autonomy.

It is **appropriate autonomy within explicit boundaries**.

29. The Development Loop as a Learning System

The agent can improve within a task because each verification cycle generates information.

Consider:

Iteration 1

Hypothesis: cache is stale.

Test

→ cache is correct.

New hypothesis:

database transaction is incomplete.

Test

→ transaction commits after invalidation.

Fix

→ move invalidation after commit.

Test

→ passes.

The system is performing a form of **empirical hypothesis testing**.

This is an important difference between code generation and agentic development.

The agent is not merely generating text.

It is interacting with an environment and updating its beliefs based on observations.

30. The Most Important Design Principle

The central lesson of Chapter 18 can now be stated precisely:

Do not optimize only for the quality of the agent's first action. Optimize for the quality of the entire trajectory toward a verified state.

This changes what we should measure.

Instead of asking:

How often does the model generate correct code?

ask:

How often does the agent reach a verified solution?

How many iterations does it require?

How much does each iteration cost?

How often does it recover from errors?

How often does it regress?

How reliably does it stop?

The relevant object is no longer the individual completion.

It is the **development trajectory**.

31. Exercise — Build an Autonomous Development Loop

Extend the coding agent built on Days 15–17.

Implement the following loop:

1. Read specification.
2. Construct task context.
3. Generate implementation plan.
4. Execute the plan.
5. Run deterministic verification.
6. Collect failures.
7. Diagnose failures.
8. Modify implementation.
9. Retest.
10. Repeat until success or stopping condition.

Add at least five verifier types:

compiler
unit tests
type checker
static analyzer
benchmark

If possible, also add:

security scanner
LLM evaluator
integration tests
browser tests

Instrument the loop.

Record:

iteration
actions
tests passed
tests failed
verification scores
tokens
cost
latency
files changed
regressions
final outcome

Then deliberately introduce bugs.

For example:

Bug 1:
Incorrect API behavior.

Bug 2:
Type error.

Bug 3:
Performance regression.

Bug 4:
Security vulnerability.

Bug 5:
Regression in an existing feature.

Observe which verifier detects each bug.

Construct a table:

Bug
↓
Verifier
↓
Feedback quality
↓
Agent response
↓
Iterations to repair

The goal is to understand that **different verifiers provide different forms of information.**

32. A More Advanced Exercise — Improve the Loop Without Changing the Model

Run the same task with the same model.

Create three systems.

System A

LLM
↓
code

System B

LLM
↓
code
↓
tests
↓
feedback
↓
LLM

System C

LLM
↓
specification
↓
context engineering
↓
plan
↓
implementation
↓
tests
↓
type checker
↓
static analysis
↓
benchmark
↓
security scan
↓
feedback
↓
replan
↓
repair
~ (loop back)

Keep the model constant.

Compare:

success rate
time to solution
token consumption
cost

number of defects
regressions
human interventions

This experiment demonstrates one of the central ideas of modern AI engineering:

A better system loop can extract substantially more capability from the same underlying model.

33. Key Takeaways

1. **Agentic development is fundamentally a feedback loop.**


```
SPEC
↓
PLAN
↓
IMPLEMENT
↓
TEST
↓
VERIFY
↓
FIX
↓
RETEST
~ (loop back)
```

2. **The goal is not perfect first-pass generation.** The goal is reliable convergence toward a verified system state.
3. **Verification is the critical feedback mechanism.** It transforms uncertain model output into evidence.
4. **Testing is only one form of verification.** Compilers, type checkers, static analyzers, benchmarks, security scanners, simulators, browser tests, and evaluators provide different forms of evidence.
5. **Different verifiers observe different dimensions of correctness.** A passing unit test does not establish performance, security, architectural compliance, or usability.
6. **Good feedback is diagnostic, not merely binary.** “Failed” is much less useful than precise evidence describing what failed, where, and under what conditions.
7. **The verifier acts like a sensor in a control system.** The agent uses its observations to choose the next action.

8. **Verification should be as independent as possible from generation.** Independent deterministic checks reduce the risk of correlated model errors.
9. **Iteration creates a new capability regime.** The agent can compensate for imperfect initial reasoning through repeated observation, diagnosis, and repair.
10. **Autonomous loops require explicit stopping conditions.** Success, iteration limits, cost limits, stalled progress, unrecoverable errors, and human-approval boundaries must all be handled.
11. **Regression control is essential.** A repair that fixes one failure while introducing another is not necessarily progress.
12. **The quality of the development trajectory matters more than the quality of an individual completion.**

$\text{Agent Capability} \approx \text{Reasoning} \times \text{Feedback Quality} \times \text{Iteration Quality}$

13. **The most powerful optimization may not be a better model.** Better specifications, context, tools, verifiers, and recovery mechanisms can substantially improve the same model's performance.
14. **Agentic development resembles closed-loop control.**



15. **The central engineering principle is simple:**

Don't make the model smarter. Make the feedback loop better.

The most important transition is from **code generation** to **verified convergence**.

A coding agent does not need to be infallible.

It needs to be embedded in a system that can reliably detect when it is wrong, determine why, make a correction, and establish that the correction actually worked.