

Contents

Specification Engineering	1
1. The Specification Engineering pipeline	2
2. What does a Specification Engineer actually do?	4
1. Intent elicitation	4
2. Behavioral modeling	4
3. Contract definition	6
3. The specification should contain invariants	6
4. Specifications should describe failure semantics	7
5. Specification Engineering includes AI-specific behavioral contracts .	8
Model behavior	8
Grounding	8
Uncertainty	8
Tool usage	8
Agent permissions	8
Stopping conditions	8
6. Specifications become generative	9
7. This makes specifications closer to executable contracts	10
Level 1 — Prose specification	10
Level 2 — Structured specification	11
Level 3 — Executable specification	11
8. Specification Engineering needs its own toolchain	11
9. AI can participate in Specification Engineering	13
10. Specifications should have traceability	13
11. Specification becomes the stable artifact in an AI-native codebase	14
12. I would define Specification Engineering this way	15
The resulting AI-native development loop	17

Specification Engineering

I would treat **Specification Engineering** as a distinct engineering discipline sitting between product intent and implementation.


The core idea is:

Specification Engineering is the systematic transformation of ambiguous human intent into precise, machine-consumable, testable, and maintainable specifications from which implementation and verification can be derived.

In traditional software engineering, this activity is distributed across product managers, architects, designers, and engineers. In an AI-native environment, it becomes much more explicit because **the specification is one of the primary inputs to the coding agent.**

1. The Specification Engineering pipeline

I envision something like:



The important change is that **the specification becomes a central engineering artifact rather than merely documentation.**

2. What does a Specification Engineer actually do?

I would break the discipline into roughly eight activities.

1. Intent elicitation

Start with something messy:

“We need a research assistant that can answer questions over our internal documents.”

The Specification Engineer extracts:


- actors
- goals
- workflows
- assumptions
- constraints
- business rules
- security requirements
- quality attributes
- unacceptable behaviors

This is partly product analysis and partly systems engineering.

2. Behavioral modeling

Turn prose into observable behavior.

For example:



Now you can define:

- states
- transitions

- events
- preconditions
- postconditions
- failure transitions

This is essentially **formal state-machine thinking applied pragmatically**.

3. Contract definition

Define what components promise to each other.

For example:

POST /documents

Input:

file: PDF
max_size: 50 MB

Output:

document_id: UUID
status: PROCESSING

Guarantees:

document_id is globally unique
caller can access only their own document

Errors:

400 malformed file
413 file too large
401 unauthenticated

The contract becomes useful simultaneously to:

- the coding agent
 - the API designer
 - the test generator
 - the reviewer
 - the monitoring system
-

3. The specification should contain invariants

This is one of the biggest differences from conventional requirements.

An **invariant** says something that must *always* remain true.

For example:

INV-001:

A user must never retrieve a document belonging to another user.

INV-002:

A deleted document must never appear in search results.

INV-003:

Every completed payment has exactly one corresponding order.

INV-004:

A retry of an idempotent operation must not create duplicate side effects.

These are extraordinarily valuable for AI agents because they constrain the solution space.

The agent can generate many implementations, but every implementation must satisfy the invariant.

4. Specifications should describe failure semantics

Traditional requirements tend to emphasize the happy path.

AI-generated systems need much stronger failure specifications.

Instead of:

“The system retrieves documents.”

specify:

Retrieval failure:

If vector search times out:

 retry once

If retry fails:

 fall back to lexical search

If lexical search fails:

 return RETRIEVAL_UNAVAILABLE

Never:

 fabricate search results

 return results from unauthorized documents

This is particularly important because an LLM-based system has **probabilistic failure modes** that don't exist in conventional deterministic software.

5. Specification Engineering includes AI-specific behavioral contracts

This is where I think the discipline becomes genuinely new.

A specification for an AI system may need to describe:

Model behavior

The model must answer only from supplied evidence.

Grounding

Every factual claim must be attributable to one or more retrieved passages.

Uncertainty

If evidence confidence < threshold:
do not provide a definitive answer.

Tool usage

The financial database must be queried before answering questions involving current account balances.

Agent permissions

The agent may:
read documents
search database

The agent may not:
delete documents
execute arbitrary SQL
modify financial records

Stopping conditions

The agent must terminate when:
answer confidence >= threshold
OR
maximum iterations = 8
OR
no additional information can be obtained.

This is essentially **behavioral specification for probabilistic components**.


6. Specifications become generative

This is perhaps the most interesting consequence.


A good specification should not merely tell an engineer what to build.

It should allow an AI system to derive engineering artifacts.

For example:



And then:



That starts to resemble a **specification** → **synthesis** → **verification loop**.

7. This makes specifications closer to executable contracts

I would distinguish three levels:

Level 1 — Prose specification

The system should respond quickly.

Weak.

Level 2 — Structured specification

```
p95 latency < 500 ms  
under 100 requests/sec  
for payloads <= 1 MB
```


Much better.

Level 3 — Executable specification

```
assert p95_latency(  
    workload="standard",  
    rps=100,  
    payload_size <= 1MB  
) < 500ms
```

Now the specification can directly participate in verification.

The long-term trajectory is:



Not every requirement needs to reach Level 4. But the important ones should.

8. Specification Engineering needs its own toolchain

I would expect a future Specification Engineering environment to look something like an IDE.

Instead of primarily editing source code, you might have:

+-----+-----+	
Product: Research Assistant	
+-----+-----+	
REQUIREMENTS	
R-001 Document ingestion	
R-002 Grounded answers	
R-003 Citation support	
BEHAVIOR	
B-014 Search workflow	
B-015 Answer workflow	
INVARIANTS	
I-001 Authorization isolation	
I-002 Citation grounding	
CONTRACTS	
API-001 /documents	
API-002 /search	
EVALS	
E-001 Retrieval accuracy	
E-002 Groundedness	
E-003 Hallucination	
AGENTS	
research-agent	
retrieval-agent	
+-----+-----+	

And the system could continuously answer:

Which requirements are implemented?

Which specifications have no tests?

Which tests fail?

Which agent changed behavior outside its specification?

Which specifications conflict?

Which requirements remain underspecified?

That last question is particularly important.

9. AI can participate in Specification Engineering

The specification itself can become an interactive artifact.

Human:

“Users should be able to share documents.”

Agent:

“With whom?”

Human:

“Other users in the organization.”

Agent:

“Can external users receive access?”

Human:

“No.”

Agent:

“Can the owner revoke access?”

Human:

“Yes.”

Agent:

“Should revocation invalidate existing sessions?”

Human:

“Yes.”

The agent is effectively performing **specification elicitation**.

It identifies ambiguity before implementation.


This is a major shift:

The AI agent should not merely generate code from specifications. It should help discover and eliminate specification ambiguity before code is generated.

10. Specifications should have traceability

I would make traceability a first-class property.

Something like:



Then a change to R-17 can automatically identify everything affected.

This gives you something close to **requirements traceability**, but much richer and much more automated.

11. Specification becomes the stable artifact in an AI-native codebase

This is perhaps the biggest conceptual shift.

Today:


Source code



system

Documentation describes the source code.

In an AI-heavy future:




The implementation may become increasingly disposable.

Agents can regenerate or substantially rewrite it.

The specification, contracts, invariants, evaluations, and architectural constraints become the **persistent representation of engineering intent**.

That suggests a very different lifecycle:




The system becomes a continuously evolving **specification–implementation–evaluation loop**.

12. I would define Specification Engineering this way

For your book, I think the strongest formulation is:


Specification Engineering is the discipline of transforming human intent into precise behavioral, structural, and operational contracts that can guide AI agents, constrain implementation, and generate automated verification.

And I would emphasize that it has **four objectives**:



That makes it fundamentally different from merely writing better requirements.

The resulting AI-native development loop



That, in my view, is the important idea: as implementation becomes increasingly automated, **the scarce engineering skill moves upward—from writing code toward precisely specifying what the code must mean and how we will know it is correct.**